


A High-Level Guide to

Boxes in Ideogram 4

How to **place**, **edit**, **prompt**, and **use** bounding boxes for layout-controlled image generation.



WHAT ARE “BOXES”?

In Ideogram 4, a **box** (bounding box) is a rectangular region you assign to a single element — a headline, a subject, a product, a logo, a background fill. Instead of describing a whole scene in one sentence and hoping the model arranges it well, you **declare a canvas and place each element by coordinate**, the way you would drag a frame in a design editor. Each box is an addressable element with its own position, its own style, and — for text — its own literal string.

Reference guide · prepared June 2026 · sources listed on the final page.

1 · The mental model

Ideogram 4 introduces **layout control** through a structured prompt. The shift is from *describing* a picture to *composing* one. You think like a designer laying out a poster: a title goes near the top, the product sits in the lower-right, a logo pins to a corner, a background fills the rest. Every one of those is a box.

- **Plain-text prompt** — still works. Best for loose, illustrative, or exploratory images where exact placement doesn't matter.
- **Structured / JSON prompt** — the box-driven mode. Best for posters, ads, UI mockups, packaging, and anything where text and elements must land in specific spots.
- **Magic Prompt** — a hosted layer that rewrites your plain-text idea into a full structured caption, so you get box-level layout without writing JSON by hand.

RULE OF THUMB

Reach for boxes the moment placement, multiple text blocks, or repeatable layout matters. For a single hero image with no strict composition, a sentence is enough.

2 · How boxes are addressed (the coordinate system)

Every box is positioned on a **normalized 0–1000 grid**. The origin (0, 0) is the **top-left** corner; x increases to the right and y increases downward. A box is written as four numbers:

```
bounding_box = [ y_min , x_min , y_max , x_max ]

# all values 0-1000, origin top-left
# y first, then x (note the order!)

[0, 0, 1000, 1000] -> the entire canvas
[0, 0, 200, 1000]  -> a full-width strip across the top
[450, 550, 850, 950] -> a block in the lower-right quadrant
```

Because coordinates are normalized, the same box description works at any output resolution (Ideogram 4 renders up to native 2K). Boxes may overlap — useful for placing text on top of a background — and the model treats each as an explicit instruction rather than a hint.

3 · Placing boxes

There are two ways to place a box, depending on where you are working.

A · On the Canvas (visual)

- Open the **Canvas / Editor** and start from a blank frame or an existing image.
- **Drag a rectangle** where the element should live — this is your box.
- Type the prompt for *that region only*; the rest of the image stays as context.
- Repeat for each element. Move or resize a box by dragging its handles.

B · In a structured prompt (precise)

Declare the canvas, then add one entry per element with its box and description. This is exact, repeatable, and easy to version-control.

```
{
  "canvas": { "width": 1024, "height": 1024 },
  "elements": [
    {
      "type": "background",
      "bounding_box": [0, 0, 1000, 1000],
      "desc": "soft studio gradient, warm neutral tones"
    },
    {
      "type": "text",
      "bounding_box": [60, 60, 220, 700],
      "text": "SUMMER SALE",
      "desc": "bold condensed sans-serif, deep navy"
    },
    {
      "type": "object",
      "bounding_box": [450, 550, 880, 950],
      "desc": "a glass bottle of cold-pressed juice, studio lit"
    }
  ]
}
```

Notice each element carries its own **bounding_box** and **desc**; text elements add a literal **text** string (covered next).

4 · Prompting a box (text vs. visual description)

A box does two jobs at once: it says *where* something goes and *what* it is. For **text** elements this splits into two distinct fields — and keeping them separate is the single biggest reason Ideogram 4 renders clean, accurate type.

Field	What it controls	Example
text	The exact literal characters to render — spelled exactly as they should appear.	"GRAND OPENING"
desc	The visual styling — font feel, weight, color, finish. No content here.	"bold serif, gold foil, slight emboss"

Because each box styles its own text independently, you can mix fonts and treatments in one image — a heavy display headline, a light subhead, and a small caption strip, each in its own box. This per-element control is what makes multi-line, multi-font in-image text reliable.

Color and palette

Layout pairs with **palette conditioning**: supply up to **16 hex colors** to constrain the whole composition, and set per-element colors inside each box's **desc** for finer control. Keep brand colors consistent by reusing the same hex values across boxes.

WRITING GOOD BOX DESCRIPTIONS

- Put the **literal words** in text, never in desc.
- Describe **style, not content** in desc.
- Size the box to the content — cramped boxes shrink or clip text.
- Leave breathing room; don't pack every region edge-to-edge.
- Reuse exact hex codes across boxes to lock brand color.

5 · Editing boxes

Editing is where the two modes diverge most — and it's important to understand the difference so you pick the right tool.

A · Editing the structured prompt (re-generation)

Change a box's coordinates, swap its desc, or rewrite its text, then run again. This is fast and precise for iterating on layout. **Caveat:** editing a JSON field *regenerates the entire image from scratch*. Even if you changed one word, the characters, background, and lighting can all shift. Use this while you are still exploring a layout.

B · Editing the pixels (localized, preserves the rest)

When you have an image you like and want a **surgical change** — swap the text on a sign, change a product label, replace one object — use the **Canvas editor / Magic Fill** (Ideogram 3.0-style Edit). You select the region (your box), prompt only that area, and the rest of the frame is preserved. This is true localized editing, not a re-roll.

You want to...	Use
Explore / restructure the whole layout	Edit the structured JSON and re-run
Reposition or resize an element quickly	Drag the box on the Canvas
Change one thing, keep everything else	Magic Fill / region edit (Edit mode)
Fix a typo in rendered text	Region-select the text, edit that box only

6 · A practical end-to-end workflow

- 1 **Sketch the layout.** Decide the major regions — headline, subject, supporting text, logo, background — before touching a prompt.
- 2 **Declare the canvas** and add a full-canvas background box first so the scene has a base.
- 3 **Place each element** with its bounding box, working back-to-front (background → subjects → text → logo on top).
- 4 **Prompt each box:** literal text + visual desc; add palette hex codes for brand color.
- 5 **Generate and review** placement. Adjust coordinates and re-run while the layout is still settling.
- 6 **Lock it in.** Once you love the composition, switch to **Magic Fill / region edit** for any final, surgical fixes so the rest is preserved.
- 7 **Export** at the resolution you need (up to native 2K).

Common pitfalls

- **Wrong coordinate order.** It's [y_min, x_min, y_max, x_max] — y comes first.
- **Content in the wrong field.** Literal words belong in text; desc is style only.
- **Expecting a JSON edit to be surgical.** Re-running regenerates everything; use region edit for localized changes.
- **Over-stuffed boxes.** Too much text in a small box clips or shrinks — size to fit.
- **Inconsistent color.** Reuse the same hex values rather than re-describing colors in words.

Quick reference

Box = one element's rectangular region. **Coords** = [y_min, x_min, y_max, x_max], 0–1000, top-left origin.

text = literal characters. **desc** = visual style. **palette** = up to 16 hex colors.

Edit JSON → re-generates all. **Magic Fill** → edits one region, keeps the rest.

Sources

- Ideogram official docs — Prompt Box & Editor (docs.ideogram.ai)
- Ideogram 4.0 technical announcement (ideogram.ai/blog/ideogram-4.0)
- Runware — Structured prompts for Ideogram 4.0
- Ideogram 4.0 prompt & layout-control guides (imagine.art, theplanettools.ai, creeta.com)